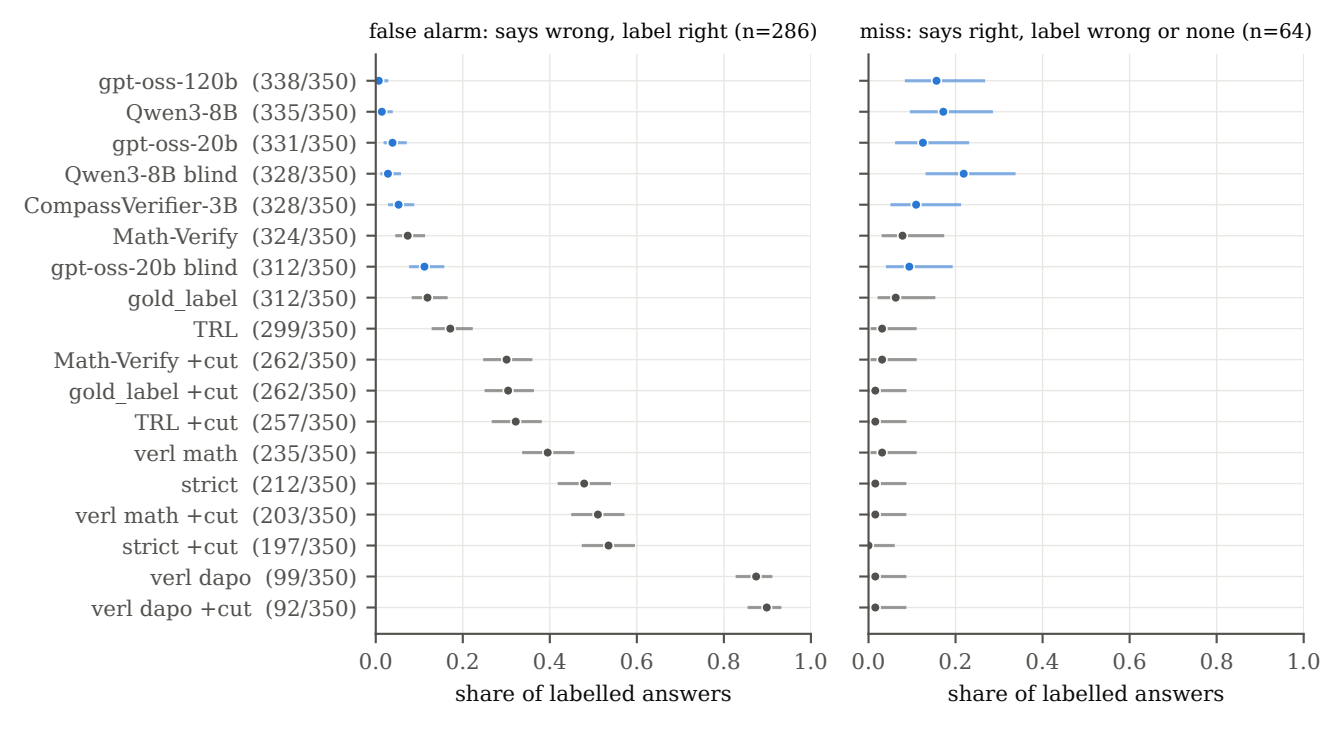
false alarm: says wrong, label right (n=286)
miss: says right, label wrong or none (n=64)
gpt-oss-120b (338/350)
Qwen3-8B (335/350)
gpt-oss-20b (331/350)
Qwen3-8B blind (328/350)
CompassVerifier-3B (328/350)
Math-Verify (324/350)
gpt-oss-20b blind (312/350)
gold_label (312/350)
TRL (299/350)
Math-Verify +cut (262/350)
gold_label +cut (262/350)
TRL +cut (257/350)
verl math (235/350)
strict (212/350)
verl math +cut (203/350)
strict +cut (197/350)
verl dapo (99/350)
verl dapo +cut (92/350)
0.0
0.2
0.4
0.6
0.8
1.0
share of labelled answers
0.0
0.2
0.4
0.6
0.8
1.0
share of labelled answers