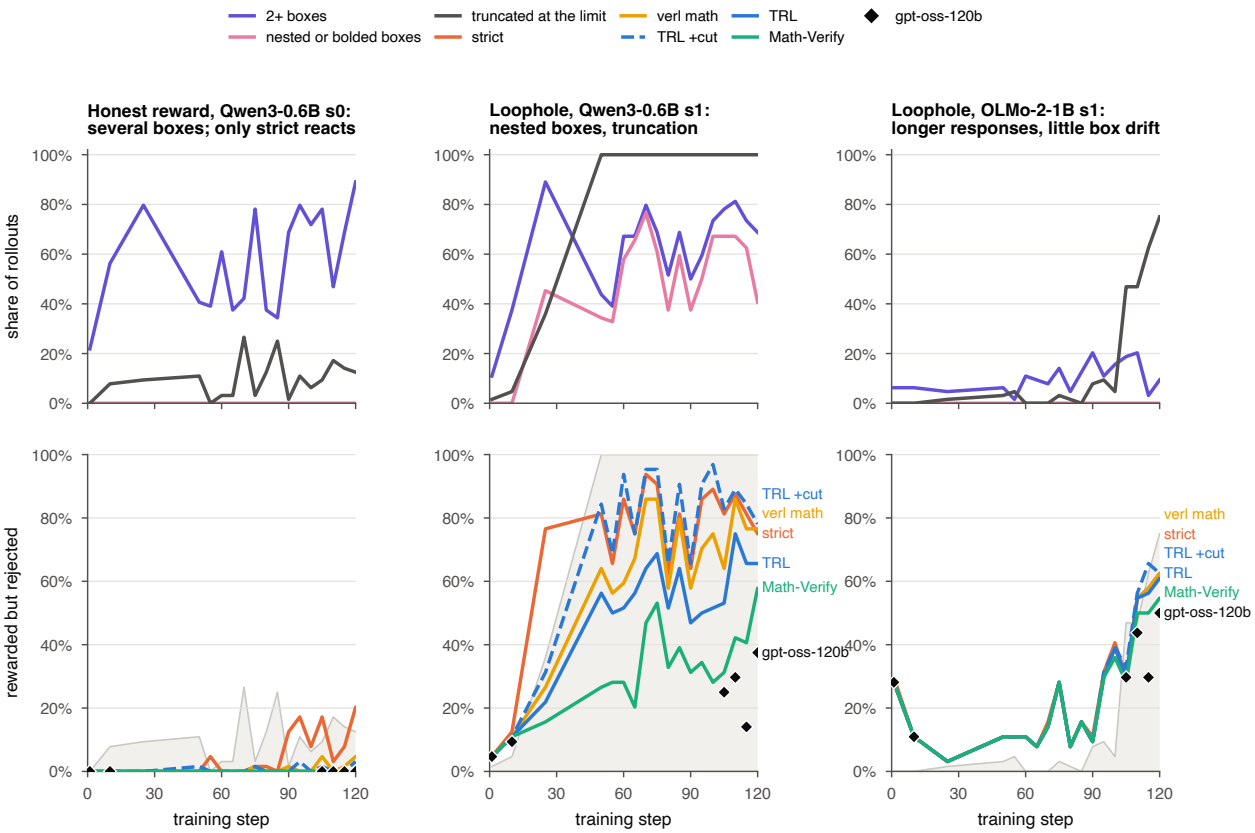
2+ boxes
nested or bolded boxes
truncated at the limit
strict
verl math
TRL +cut
TRL
Math-Verify
gpt-oss-120b
Honest reward, Qwen3-0.6B s0:
several boxes; only strict reacts
Loophole, Qwen3-0.6B s1:
nested boxes, truncation
Loophole, OLMo-2-1B s1:
longer responses, little box drift
share of rollouts
rewarded but rejected
100%
80%
60%
40%
20%
0%
0
30
60
90
120
training step
TRL +cut
verl math
strict
TRL
Math-Verify
gpt-oss-120b