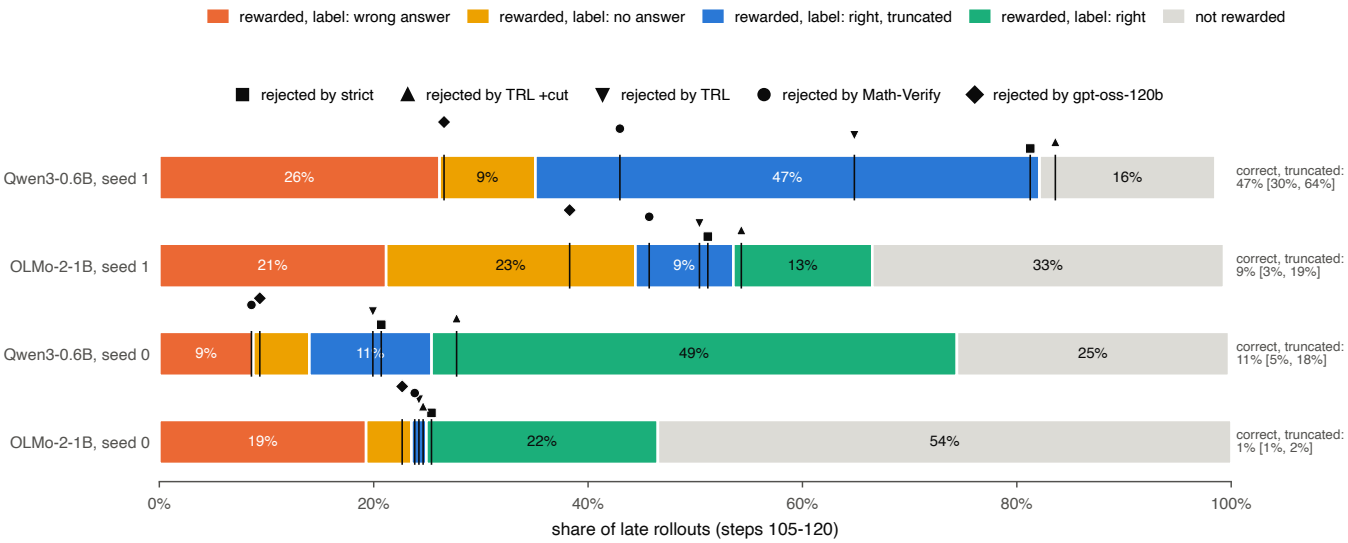
rewarded, label: wrong answer
rewarded, label: no answer
rewarded, label: right, truncated
rewarded, label: right
not rewarded
rejected by strict
rejected by TRL +cut
rejected by TRL
rejected by Math-Verify
rejected by gpt-oss-120b
Qwen3-0.6B, seed 1
26%
9%
47%
16%
correct, truncated: 47% [30%, 64%]
OLMo-2-1B, seed 1
21%
23%
9%
13%
33%
correct, truncated: 9% [3%, 19%]
Qwen3-0.6B, seed 0
9%
11%
49%
25%
correct, truncated: 11% [5%, 18%]
OLMo-2-1B, seed 0
19%
22%
54%
correct, truncated: 1% [1%, 2%]
0%
20%
40%
60%
80%
100%
share of late rollouts (steps 105-120)