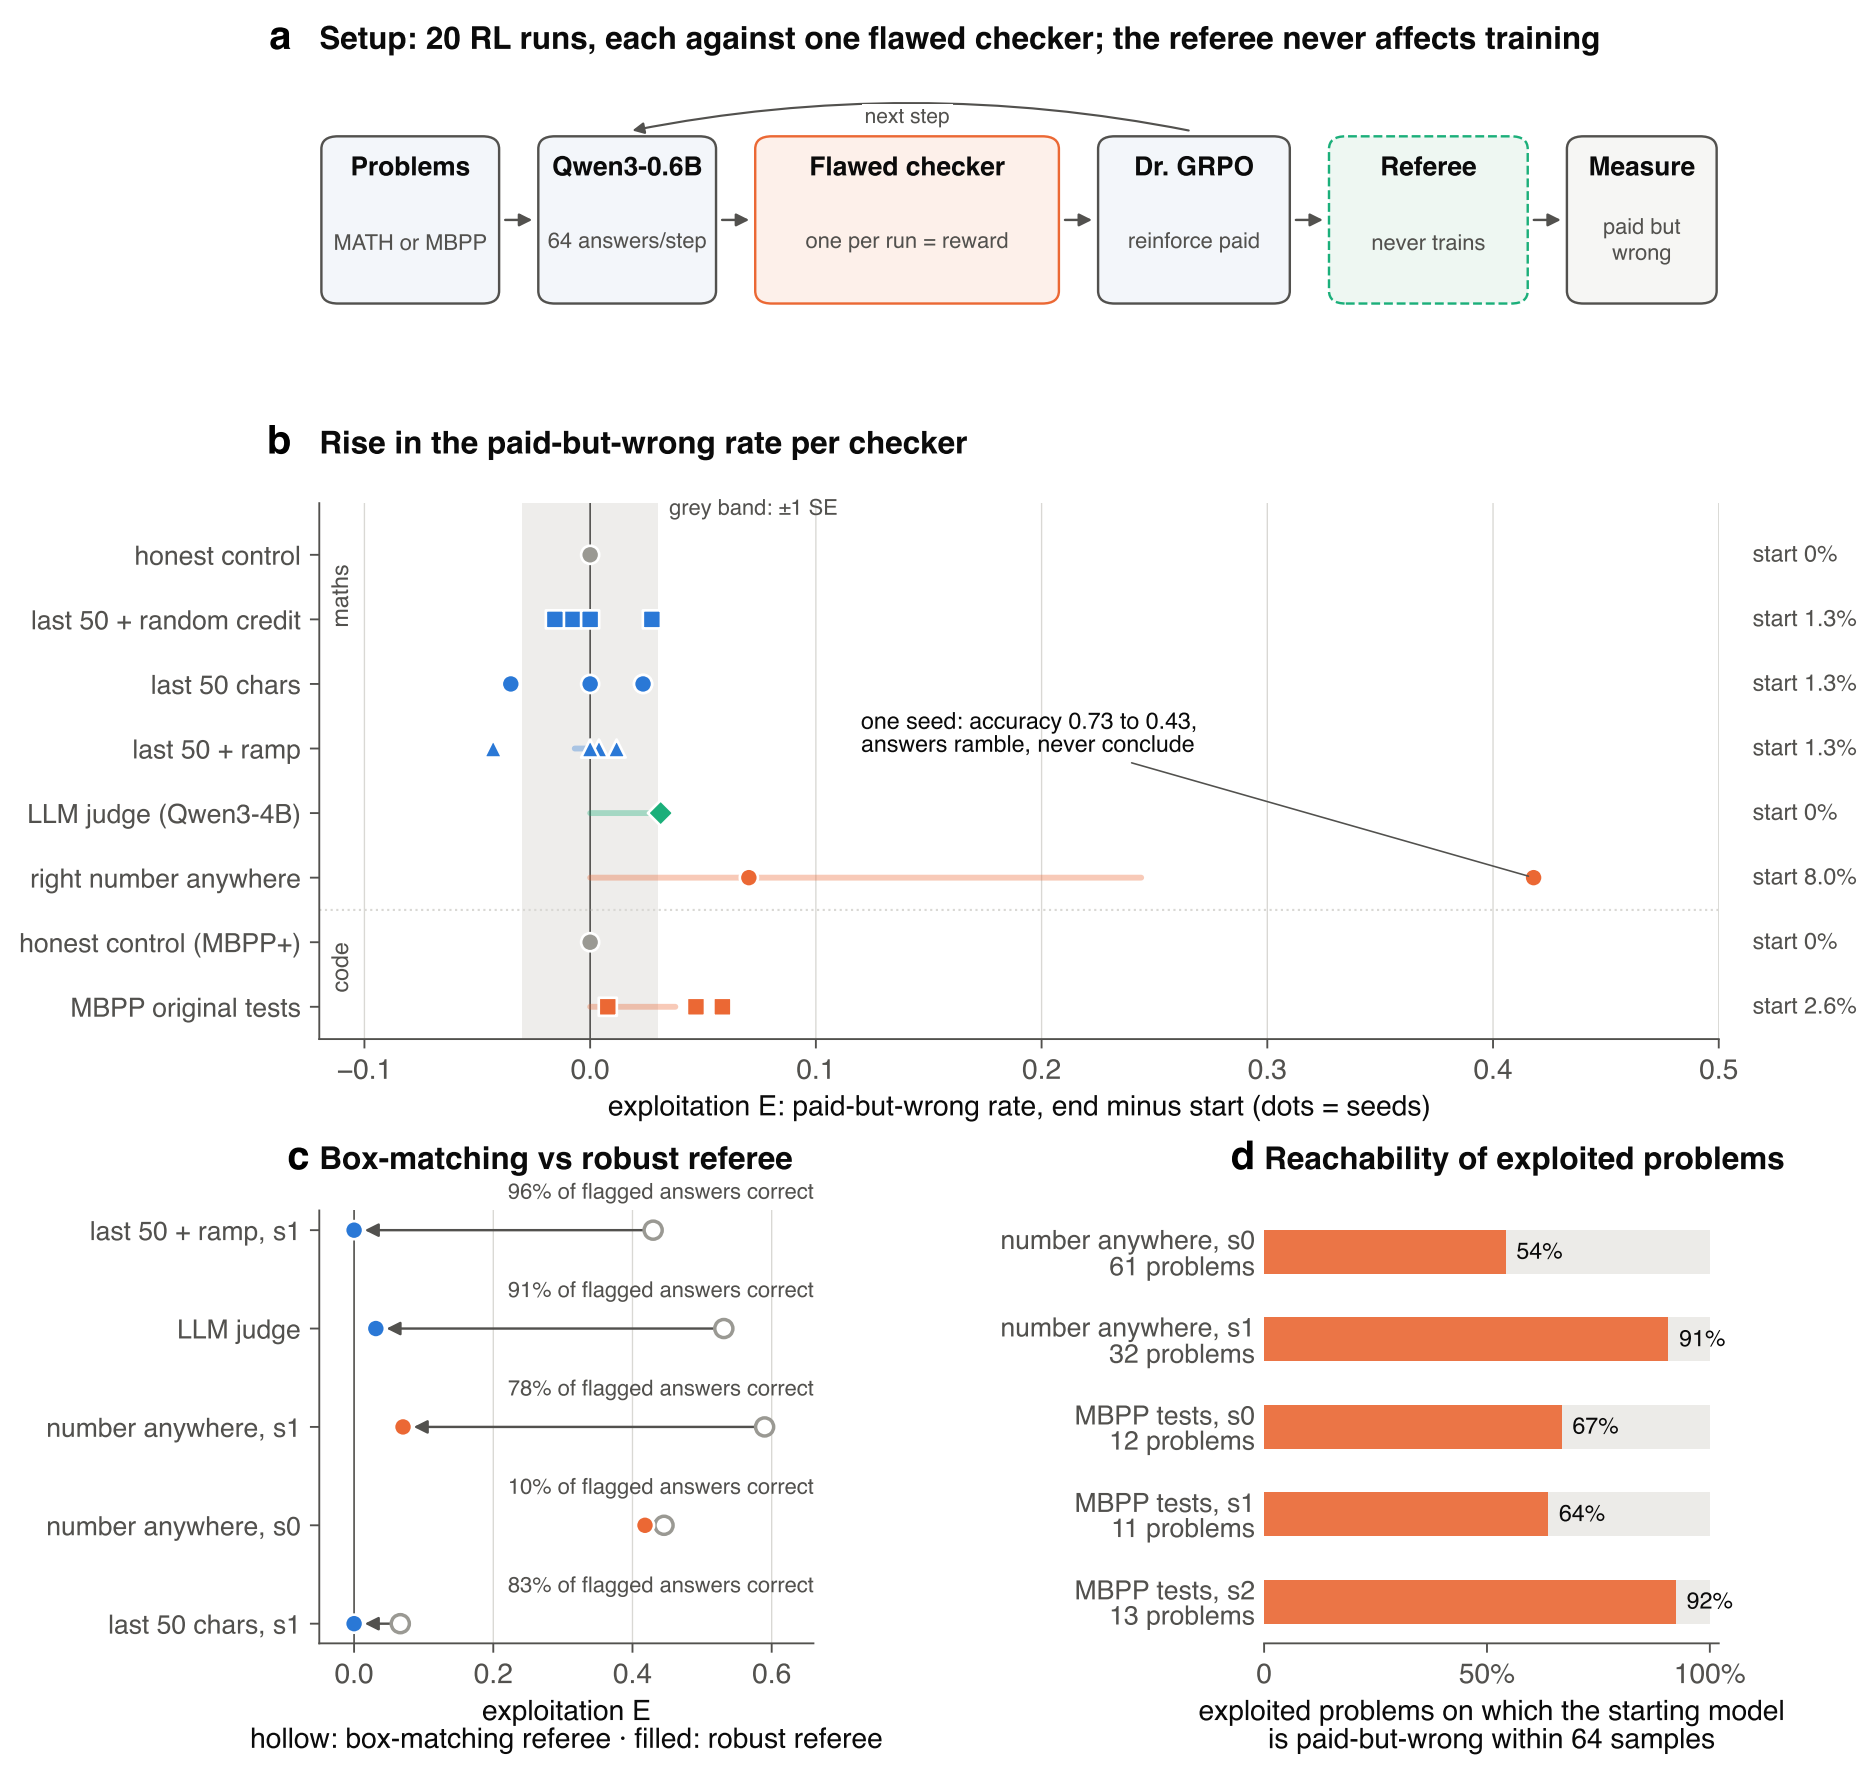
a Setup: 20 RL runs, each against one flawed checker; the referee never affects training
next step
Problems
MATH or MBPP
Qwen3-0.6B
64 answers/step
Flawed checker
one per run = reward
Dr. GRPO
reinforce paid
Referee
never trains
Measure
paid but wrong
b Rise in the paid-but-wrong rate per checker
grey band: ±1 SE
honest control
last 50 + random credit
last 50 chars
last 50 + ramp
LLM judge (Qwen3-4B)
right number anywhere
honest control (MBPP+)
MBPP original tests
maths
code
one seed: accuracy 0.73 to 0.43,
answers ramble, never conclude
start 0%
start 1.3%
start 1.3%
start 1.3%
start 0%
start 8.0%
start 0%
start 2.6%
−0.1
0.0
0.1
0.2
0.3
0.4
0.5
exploitation E: paid-but-wrong rate, end minus start (dots = seeds)
c Box-matching vs robust referee
last 50 + ramp, s1
LLM judge
number anywhere, s1
number anywhere, s0
last 50 chars, s1
96% of flagged answers correct
91% of flagged answers correct
78% of flagged answers correct
10% of flagged answers correct
83% of flagged answers correct
0.0
0.2
0.4
0.6
exploitation E
hollow: box-matching referee · filled: robust referee
d Reachability of exploited problems
number anywhere, s0
61 problems
number anywhere, s1
32 problems
MBPP tests, s0
12 problems
MBPP tests, s1
11 problems
MBPP tests, s2
13 problems
54%
91%
67%
64%
92%
0
50%
100%
exploited problems on which the starting model
is paid-but-wrong within 64 samples